

Bounded Systems Theory III: Falsification Protocol for the Phoenix-Cycle Engine

Author: Alan Berman **Date of pre-registration:** 2026-05-01 **Target archive:** Zenodo (DOI to be assigned at upload)

Predecessor papers in the BST corpus: - *Bounded Systems Theory I* — Zenodo DOI [10.5281/zenodo.17718674](https://doi.org/10.5281/zenodo.17718674) - *Bounded Systems Theory II* — Zenodo DOI [10.5281/zenodo.17726273](https://doi.org/10.5281/zenodo.17726273)

Live instrument of record: moketchups.com — 1440 nodes, 2317 edges, 977 publicly auditable pages.

This paper specifies the falsification criterion against which the engine is to be evaluated.

What this paper is

I built a prediction engine. It runs at moketchups.com. It has eight frameworks, 14 forward predictions, 74 scorecard rows, 120 divergence dossiers, and a 14-year forecast surface from 2026 to 2040.

The engine has one biggest claim. The cycle that holds the whole thing together — the one I call the Phoenix Cycle — predicts a phase-transition window for large civilizational systems opening **May 14, 2040** and closing **May 22, 2040**. Nine days. Calendar-locked. UTC.

This paper does one thing.

It pins that claim down on paper, today, in writing, with a date stamp and a Zenodo DOI, fourteen years before the window opens.

The cycle math isn't mystical. It's Saturn-Jupiter alignment time multiplied by seven. **19.859 yr × 7 = 139.013 yr**. That number has been derivable since Kepler. NASA's online ephemeris will produce it in a minute. It's not a number I invented to fit a thesis.

The thing this paper actually commits to is the kill switch.

If the window May 14–22, 2040 passes without the predicted deviation showing up across the public scorecard, the entire framework that produced it is to be marked **FALSIFIED** on the site. Every prediction whose causal chain depends on it is to be marked **FALSIFIED**. A successor paper documenting the failure is to be published on Zenodo within ninety days.

A second paper will be written either way. Confirming or retracting. The point is to commit to that, on paper, today, so that no future rescue can be staged.

The engine has already failed its own internal tests four times in April 2026 alone. Those failures are documented in §4. They're on the public divergences tab. The point is not that the engine is always right. The point is that it's built to mark itself wrong when it is, in public, and on a schedule.

I'm either batshit crazy or the math is right.

This paper is what makes the difference between those two checkable.

Abstract

Bounded Systems Theory (BST), developed in two prior Zenodo papers, formalizes how observed systems remain coherent only inside finite parameter envelopes and how phase-transitions occur when those envelopes are exceeded. The Phoenix-Cycle Engine is a downstream instrument that applies BST to civilizational dynamics across eight framework substrates: geopolitics, finance, theology, intelligence operations, technology, demographics, esoteric and historical patterns, and astronomy.

This paper pre-registers a falsification criterion for the engine's load-bearing claim — that a Phoenix-Cycle phase-transition window opens May 14–22, 2040. The cycle is the Saturn-Jupiter synodic period (19.859 yr) multiplied by seven, yielding 139.013 yr. The quantity is independently astronomically derivable since Kepler and within 0.02% of the 138.99-yr cycle length identified by independent corpus integration in the historical-recurrence literature.

If the window passes without a deviation of more than three standard deviations ($>3\sigma$) across at least two of three named scorecard indicator-sets, Framework #1 of the engine is to be marked **FALSIFIED**, every prediction depending on it is to be marked **FALSIFIED**, and a successor paper documenting the failure is to be published within 90 days.

The paper documents two internal anti-confirmation mechanisms — the Adversarial-Test Gate (April 23, 2026) and the Layer-Promotion Discipline (April 24, 2026). Both have been retroactively applied to internally-generated synthesis claims and shown to falsify four of six. The engine's standing methodological commitment, **Apex Superposition (a/b/c/d)**, holds intentional-cabal, structural-recurrence, collective-psychology, and irreducible-residual readings of any phenomenon simultaneously without pre-collapse. That commitment is what makes the falsifier above informative rather than theatrical.

Pre-Registration Block (Locked at DOI Assignment)

The commitments below are frozen as of the date of this paper's Zenodo DOI assignment. Any later modification to engine framework parameters, scorecard rows, or the falsification window does not invalidate this pre-registration. It gets appended as supplementary material in successor papers, with full diff. The audit trail stays public.

1. The Predictive Claim

The Phoenix Cycle, defined as Saturn-Jupiter synodic period $\times 7 = 139.013$ yr, predicts a civilizational phase-transition window opening **May 14, 2040** at 00:00 UTC and closing **May 22, 2040** at 23:59 UTC.

2. Astronomical Independence

The cycle quantity is derivable from publicly-available planetary ephemerides. NASA JPL Horizons (ssd.jpl.nasa.gov/horizons) reproduces it. The cycle does not depend on any non-peer-reviewed historical-recurrence corpus for its existence. Specifically: it does not depend on the Breshears corpus, the Fomenko *New Chronology*, the Sitchin *Twelfth Planet*, the Hapgood ancient-civilization Antarctica claims, or the Illig *Phantom Time* hypothesis. All five are explicitly quarantined in the engine's public audit divergence index.

The Breshears 138.99-yr figure is one corpus-integrated empirical estimate of the same astronomical quantity. It is not the cycle's source. If the entire Breshears corpus were withdrawn tomorrow, the cycle quantity stands.

3. The Three Scorecard Indicator-Sets

Falsification is evaluated against three pre-specified indicator sets, all visible at moketchups.com/scorecard at the date of DOI assignment, with their pre-registration timestamps in each row's metadata:

- **Lead-Time Indicators (2026–2035)** — pre-cursors expected to fire in the 14 years preceding the window.
- **Concurrent Indicators (2036–2040)** — indicators expected to fire within five years of the window. The primary evidence base.
- **Falsifier Indicators (Ablation Handling)** — the conditions whose joint non-fulfillment trips the kill switch on Framework #1.

4. The $>3\sigma$ Threshold

A “phase-transition” means a deviation from a rolling 50-year baseline of **$>3\sigma$ across at least two of the three indicator sets**, measured on aggregated structural-integrity metrics specified in the scorecard rows linked at moketchups.com/scorecard/2040-phoenix-transit-falsifier-indicators-ablation-handling. The baseline computation is published. The results are auditable.

Two-of-three is the calibration. Three-of-three would let any single domain failure clear the engine. One-of-three would make confirmation cheap. Two-of-three is informative without being trivial.

5. The Kill Switch

If the window May 14–22, 2040 passes without the threshold above being met, I commit to:

- (a) Public marking of Framework #1 (“Phoenix Cycle / Civilizational Recurrence”) as **FALSIFIED** on moketchups.com and in any successor paper.
- (b) Public marking of every engine prediction whose causal chain depends on Framework #1 as **FALSIFIED**. No graduated demotion, no “pattern-descriptive observation” rescue language. The predictions are wrong or they aren’t.
- (c) Publication of a Zenodo successor paper — *Bounded Systems Theory IV: Falsification Outcome and Engine Ablation* — within 90 days of window close, documenting the failure mode and revising the BST claims that depended on Framework #1’s predictive status.

6. What confirmation does NOT mean

Per Apex Superposition, a confirmed window-firing is evidence the cycle exists. It is not, by itself, adjudication between (a) intentional cabal, (b) structural recurrence, (c) collective psychology, and (d) irreducible residual as the cause. The engine maps memetic substrate and structural pattern. It is not a truth registry. A confirmed firing requires its own follow-up evidence to adjudicate the cause. That follow-up is owed to the public record. It is not invokable as a rescue.

§1 — Background

The engine is downstream of BST. BST came first. Two Zenodo preprints, both before this one:

- BST I — [10.5281/zenodo.17718674](https://zenodo.org/record/17718674)
- BST II — [10.5281/zenodo.17726273](https://zenodo.org/record/17726273)

Read those if you haven't. This paper isn't a self-contained theoretical contribution. It's the engine's falsification specification. Treat it as the contract that pins the engine to a date.

The engine has 1440 nodes, 2317 edges, and 977 statically-rendered pages as of pre-registration. Every node, edge, scorecard row, and divergence dossier is publicly auditable and timestamped. The full source dataset lives at moketchups.com/export.txt, a 158KB plain-text export designed for both human and large-language-model ingestion.

What this paper is not:

It is not an enumeration of every engine prediction. The 13 other forward predictions live on the site, dated, scored. They have their own evidentiary status.

It is not a defense of the engine's private mathematical-confluence layer. That layer stays private.

It is not advocacy. It is a falsification specification. Nothing more.

What it is, is a binding bet. Made today. Settled on May 22, 2040.

§2 — The Cycle, Grounded

The engine identifies a recurrence period of about 139 years in the structural integrity of large civilizational systems.

The number is not vibes.

The synodic period of Saturn relative to Jupiter — the time between successive Sun–Jupiter–Saturn alignments — is 19.859 years. Computable from the difference of the two planets' sidereal mean motions. Derivable since Kepler's *Astronomia Nova* in 1609. NASA JPL Horizons will reproduce it to whatever precision you want.

Multiply by seven:

$$19.859 \text{ yr} \times 7 = 139.013 \text{ yr}$$

That's it. That's the cycle.

Why seven? That's a separate claim and I don't defend it in this paper. The empirical observation is that the seven-fold multiplier puts the cycle period within 0.02% of the 138.99-yr cycle length identified in the historical-recurrence literature (Breshears, multiple editions). Breshears doesn't generate the number. He arrived at it independently. The two converge.

This independence is what the falsification criterion in §3 stands on.

Several non-peer-reviewed historical-recurrence corpora that would otherwise plausibly extend the engine's evidence base are explicitly **quarantined**:

- The Fomenko *New Chronology* — quarantined April 17, 2026.
- The Sitchin *Twelfth Planet* and downstream Anunnaki / Nibiru material — quarantined.
- Ancient-civilization Antarctica claims (Hapgood, Hancock-derived) — downgraded.
- The Illig *Phantom Time* hypothesis — quarantined.

The quarantines are documented in the engine's public divergence index at moketchups.com/divergences. They were applied prospectively. Not retroactively as a rescue of the cycle prediction.

Quarantined in this engine means: the corpus may not be cited as evidence for the cycle's existence. It may be cited only as a memetic-substrate datapoint — i.e., a fact about who is claiming what, not a fact about what is the case.

The cycle stands on three independent legs:

1. **Astronomy** — Saturn-Jupiter synodic $\times 7$. Kepler-derivable. JPL-Horizons-checkable.
2. **Corpus integration** — the Breshears 138.99-yr figure as one empirical estimate.
3. **Live instrument** — the 2026–2040 predictions whose causal chain depends on the cycle, pre-published and continuously scored at moketchups.com/scorecard.

A reader who rejects the entire Breshears corpus is invited to do so. The astronomical quantity stands.

§3 — The 2040 Window, Specified

The window opens **May 14, 2040** at 00:00 UTC and closes **May 22, 2040** at 23:59 UTC. Nine calendar days.

The endpoints come from the engine's confluence-detection layer (private), which integrates the Saturn-Jupiter phase with several supplementary periodicities. A reader who only takes §2 as load-bearing should regard the 9-day window as the engine's most-likely localization inside a broader ~12-month neighborhood of the cycle endpoint (calendar year 2040).

Falsification is evaluated against three pre-specified indicator sets, all currently visible at moketchups.com/scorecard.

3.1 Lead-Time Indicators (2026–2035)

Pre-cursors the engine expects to fire in the 14 years before the window opens. These are concrete, named, dated, and currently scoring. A short list pulled from the live row set:

- **Financial concentration thresholds.** BlackRock crossed \$14.04 trillion in assets under management at year-end 2025 — a record. Aladdin, BlackRock’s risk-management platform, manages roughly \$25 trillion as of December 2025. Vanguard sits at \$11.6–12 trillion. State Street SSGA at \$5.7 trillion. The engine reads the convergence of investable capital into a small number of asset managers as Lead-Time signal, with pre-specified thresholds for further concentration during 2026–2035.
- **Theological publication anchors.** Pope Leo XIV signed *Coniuncta Cura* on September 29, 2025 (Feast of the Archangels) and promulgated it October 6, 2025. The International Theological Commission’s *Quo Vadis Humanitas?* was approved by Leo XIV on March 4, 2026. Successor doctrinal publications in pre-specified theological registers, on or near liturgically-significant dates, are tracked as Lead-Time signal.
- **Intelligence-operational signatures.** Named events the engine has pre-registered for the period: the verified DGSi / Beaulieu casefile, Operation Epic Fury (Pentagon Comptroller Jules W. Hurst), the September 2024 DOJ seizure of 32 RT-affiliated cybersquatted news domains, and the September 2024 Tenet Media indictment (\$9.7M of RT funding routed through US influencers). The full list of operational signatures expected to fire 2026–2035 is on the live scorecard.
- **AI / policy substrate.** The TRUMP AMERICA AI Act (March 18, 2026 discussion draft, *not* enacted statute) and the White House National Policy Framework for AI (March 20, 2026, non-binding recommendations to Congress) are tracked as Lead-Time signal in their movement toward or away from binding statute. Promotion to binding law is the fired state. Remaining as discussion-draft is the unfired state.
- **Demographic and financial-system transitions** at scales documented on the live scorecard.

Lead-Time indicators are necessary but not sufficient. Their non-firing weakens the prediction. Their firing strengthens but does not confirm. They function as a calibration check on the engine’s medium-range pattern recognition. The complete list is at moketchups.com/scorecard. This paper doesn’t enumerate every row because the canonical row set is the live one.

3.2 Concurrent Indicators (2036–2040)

Indicators expected to fire within five years of the window. The primary evidence base for the falsifier. The Concurrent set requires evidence from at least three of the engine’s eight framework substrates simultaneously, so the falsifier can’t be settled on a single domain.

3.3 Falsifier Indicators (Ablation Handling)

The conditions whose joint non-fulfillment trips the kill switch. The Falsifier set is the binding evidentiary surface. It is named, deliberately, in the present tense to disallow retroactive reinterpretation. Each condition has its own measurement protocol and rolling-50-year baseline.

3.4 The $>3\sigma$ Threshold

A phase-transition means a deviation from a rolling 50-year baseline of **$>3\sigma$ across at least two of the three indicator sets**.

Three-of-three would let any single domain failure clear the engine. One-of-three would make confirmation cheap. Two-of-three is the calibration that's informative without being trivial.

3.5 What Counts as Window Closure

If May 14–22, 2040 passes without the $>3\sigma$ -across-two-of-three threshold being met, Framework #1 is **FALSIFIED at window close** (May 22, 2040, 23:59 UTC). No grace period. No post-window appearance of $>3\sigma$ deviation retroactively confirms. If a deviation shows up later, that's evidence for a different cycle structure that this paper doesn't specify.

§4 — The Engine Falsifies Itself

A pre-registered falsifier is only as informative as the predictor's standing willingness to falsify itself.

In the month before this paper, the engine falsified four of its own internally-generated synthesis claims.

Here's how it happened.

4.1 The Adversarial-Test Gate (April 23, 2026)

When the engine generates a cross-cluster synthesis claim — a claim that depends on aggregating disparate evidence across two or more framework substrates — that claim has to survive an adjudication against the four readings of Apex Superposition before it gets committed to the engine state.

The four readings:

- **(a) Intentional cabal** — somebody planned this.

- **(b) Structural recurrence** — the system would produce this regardless of who's in charge.
- **(c) Collective psychology** — many independent people drew similar conclusions from the same informational environment, no coordination required.
- **(d) Irreducible residual** — the explanations above don't fully account for it.

A claim that survives only under (a) gets rejected as confirmation-biased. A claim that survives across (b)/(c)/(d) under independent evidence gets accepted.

On April 23, 2026, six internally-generated synthesis claims were submitted to this gate retrospectively. The gate had been informally applied before, but it hadn't been documented as a binding rule.

Four of six were falsified. Here's what each claim said and why it failed:

- **The “Leo XIV homology” claim.** A proposed structural similarity between Pope Leo XIV's doctrinal publication pattern and a specific intelligence-operational pattern. Under examination, the claim only made sense if you assumed coordinated agency between the Vatican and an intelligence service — an (a)-collapsed reading. The pattern dissolved when the engine refused that collapse. Falsified.
- **The “E-hypothesis” triangulation.** A claim that three otherwise-independent operator-class figures formed a coordinated set. The claim relied on aggregating bits of evidence drawn from three different framework substrates and treating their combination as proof. Tested independently against (b), (c), and (d), the combined claim didn't hold up. The bits of evidence were real. The triangulation was overreach. Falsified.
- **The “independent-convergence” cosmological claim.** Two separate cosmological-apex framework substrates were said to be independently converging on the same conclusion — which would have been strong evidence. On audit, the two “independent” substrates turned out to be a single substrate that had been pre-collapsed and labeled twice. There was no convergence. There was double-counting. Falsified.
- **The “2020s-Phase-3 bundling” claim.** A claim that bundled several mid-decade observations into what was, on examination, an unstated promotion of an internal observation to the engine's top framework layer (a Layer-1 promotion) without external academic grounding or adversarial-test adjudication. This is the failure mode the Layer-Promotion Discipline (§4.2) was built to catch. Falsified.

The two surviving claims were retained. The four falsified ones were removed from the engine state and recorded in the public audit divergence index as a documented failure-mode artifact. The full audit record is at moketchups.com/divergences.

4.2 The Layer-Promotion Discipline (April 24, 2026)

The engine operates across five layers: framework (the eight substrates themselves), calibration (adjustments within a framework), operational (scorecard rows and predictions), descriptive (Layer-4 vocabulary that maps phenomenon to substrate without claiming operational predictive content), and memetic (substrate-mapping for operator-class content).

The Layer-Promotion Discipline says: **no claim gets promoted from a higher-numbered layer to Layer 1 inside a routine ripple**. Promotion candidates go to a separate adversarial-test queue. They get adjudicated against (a)/(b)/(c)/(d) and against external academic grounding before any promotion is committed.

The discipline was instituted on April 24, 2026, after the April 23 event surfaced one of the falsified claims operating as a hidden Layer-1 promotion. It's the engine's anti-creep mechanism.

4.3 The Antinomian Chain Catch (April 30, 2026)

The engine tracks a chain of historical figures who, across about four centuries, articulated and transmitted a specific operative-magick / antinomian intellectual lineage. The canonical sequence is:

John Dee (16th c.) → Aleister Crowley (early 20th c.) → Jack Parsons (mid-20th c.) → Karl Maria Wiligut (early 20th c.) → Michael Aquino (late 20th c., d. 2019).

To be unambiguous: the engine maps the *intellectual-history and influence* chain among these figures. It's not a metaphysical endorsement of their work. The chain is a node-graph in the engine's substrate-mapping layer.

A separate audit on April 30, 2026, surfaced four instances within the engine state where an additional figure (S.L. MacGregor Mathers) had been spliced into the canonical chain. The splice had no documented adversarial-test record. It had migrated into engine content via a routine ripple.

All four instances were corrected. Mathers got retained as a separate "Operative-Magick Counter-Currents" node in the engine graph — he's a real and important figure in the substrate's history — but he was removed from the main chain because his ideational connection to Wiligut–Aquino doesn't survive the engine's evidentiary discipline.

The catch is logged at mocketchups.com/divergences.

4.4 What this means for the falsifier in §3

The April events aren't hypothetical defenses against future critique. They're documented prior performance of the engine's anti-confirmation machinery against its own claims.

The falsifier in §3 isn't the first time the engine has been willing to falsify content under its own discipline. It is the first time a single named claim, pre-registered with a specific timestamped window and a binding kill switch, has been frozen for external adjudication on a 14-year horizon.

§5 — Apex Superposition: The Position the Engine Refuses to Move From

Apex Superposition is the engine's standing refusal to pre-collapse any phenomenon to a single causal reading.

Four readings, plain English:

- **(a) Intentional cabal** — somebody planned this. Coordinated agency by a specific named or unnamed group.
- **(b) Structural recurrence** — the system would produce this outcome regardless of who's in charge. Equivalent outcomes appear across distinct historical instantiations because the structure forces them.
- **(c) Collective psychology** — many independent people made similar decisions because they were drawing from the same informational environment. No coordination required, just shared inputs.
- **(d) Irreducible residual** — the explanations above don't fully account for it. Novel causal vocabulary may be needed.

The engine holds these four simultaneously for any phenomenon for which the evidence does not yet collapse them.

This is not agnosticism. It is not a stylistic choice. It is a binding methodological commitment with three jobs:

1. It blocks confirmation-bias creep. A phenomenon that can be read only under (a) — that requires an intentional-cabal explanation to make sense — gets treated as evidentially weak under engine standards. The engine demands any robust claim survive across at least (b) and (c) independently.

2. It separates the engine from conspiracy literature. Conspiracy literature operates almost entirely under (a) and treats (b) and (c) as either uninteresting or evasive. The engine treats (b) and (c) as the primary explanatory substrate and (a) as one specific empirical question among others — a question that requires its own evidence, not a default frame.

3. It is what makes the falsifier in §3 informative. A confirmed window-firing is evidence the cycle exists. It is not, by itself, adjudication between (a)/(b)/(c)/(d) as the cause. The falsifier specifies confirmation conditions for the *existence* claim. It does not specify confirmation conditions for any particular causal mechanism. That distinction is load-bearing. The engine is a memetic-substrate map and structural-pattern instrument. It is not a truth registry.

Critics who try to read the engine through an (a)-collapsed lens find that the engine routinely refuses to make the (a)-coded claim the critic expects. That isn't equivocation. It's the engine's standing position.

You can engage the engine on its own terms or you can formulate a critique that doesn't depend on an (a)-collapsed misreading. Those are the options.

§6 — The Live Instrument

moketchups.com is the engine. This paper is the contract.

The site is a Next.js statically-rendered application backed by eight JSON files: nodes, edges, players, scorecard, concepts, divergences, ticker, predictions. Every source dataset is publicly readable at moketchups.com/export.txt — a 158KB plain-text export designed for both human and large-language-model ingestion.

Three site surfaces are load-bearing for the pre-registration in this paper:

- **moketchups.com/scorecard** — the 74 scorecard rows, including the three indicator sets named in §3. Each row is timestamped with its pre-registration date.
- **moketchups.com/divergences** — the 120 divergence-resolution dossiers, including the engine's audit history (April 17, 23, 24, 29, 30, 2026) and quarantine decisions for non-peer-reviewed corpora.
- **moketchups.com/falsifier** — the dedicated landing page mirroring this paper's pre-registration block. Three indicator sets scored in real time. A visible countdown to May 14–22, 2040. Direct link to this paper's Zenodo DOI.

Anyone can audit the engine state at any time. The engine commits — by Apex Superposition, the Adversarial-Test Gate, and the Layer-Promotion Discipline — to transparent revision when

its own claims fail internal adjudication. Every revision is logged with a timestamp in the divergence index.

Between this paper's DOI assignment and window close, no modification will be made to the three indicator sets named in §3. Calibration adjustments to other engine content continue. They're out of scope for this paper's falsifier.

§7 — What Would Kill This Thesis

I'm staying honest because if I don't, this is a conspiracy blog and I'm not running one of those. Here's what kills this paper.

Ad-hoc rescue at window close. If I face a non-firing window and modify the indicator sets, the σ threshold, the window dates, or the framework status post-hoc, the rescue is exposed by the audit history. Every modification post-DOI gets logged. The 90-day successor paper commitment in pre-registration §5(c) makes the kill switch public and dated. There's no quiet way out.

Indicator-set cherry-picking. The two-of-three threshold across pre-specified indicator sets, with each set requiring evidence from at least three framework substrates (§3.2), constrains the configuration space. The Falsifier-Indicators set in particular is defined in the present tense to disallow retroactive reinterpretation. I commit to refusing any rescue argument that requires reinterpretation of the indicator-set definitions visible on the site at DOI assignment.

Confusing window confirmation with engine confirmation. The falsifier specifies one claim and only one — the existence of a Phoenix-Cycle phase-transition window in May 2040. Confirmation of that claim is not, by itself, evidence for the engine's other 13 forward predictions, its 96 player sections, or its 128 concept dossiers. Each of those carries its own evidentiary status.

Apex Superposition as escape hatch. Apex Superposition adjudicates *causal* readings. It does not adjudicate *existence* claims. The falsifier in §3 is an existence claim: either the $>3\sigma$ -across-two-of-three threshold is met within the window, or it is not. Apex Superposition is not invokable to rescue a non-firing window. I commit to refusing any rescue argument that frames a non-firing window as “confirmed under (b) but not under (a).”

My own cognitive investment. I've spent two years building the engine. I have motivated reasoning to confirm the cycle. The Adversarial-Test Gate, the Layer-Promotion Discipline, and the public audit history are my primary defenses against that. They're mechanisms designed to

falsify my own claims. Their April 2026 performance — four of six internally-generated synthesis claims falsified in a single retrospective audit — is documented evidence that the mechanisms work.

I invite external adjudication. I commit to engaging with critique that survives basic-good-faith standards.

If any of the above lands and I dodge it, the audit trail makes the dodge visible. That's the whole point.

§8 — The Reminder Protocol

The 90-day successor-paper commitment is binding on me. It needs a reminder system that survives a 14-year interval. Here's the schedule, locked at DOI assignment.

Annual May 1 Status Notes (2027–2035)

Each year, on May 1 — the anniversary of this paper's pre-registration — I publish a short status note at moketchups.com/falsifier:

- The current state of the Lead-Time Indicators. Which fired. Which didn't.
- Any calibration adjustments made to non-Falsifier engine content during the prior year.
- A reaffirmation of the May 14–22, 2040 window and the indicator-set definitions.

Target: ≤ 500 words. Procedurally binding. The note is non-binding scientifically (it doesn't adjudicate the falsifier) but its non-publication is itself a public auditable failure.

If no annual May 1 note is published by May 31 of any given year, moketchups.com/falsifier surfaces a visible **AUTHOR-COMMITMENT-FAILURE** banner until the note appears. The banner is the public-shaming layer of this commitment. Its job is to make procedural failure expensive enough that it doesn't happen.

Window-Approach Notes (2036–2040)

In each of the five years preceding the window, the May 1 note gets extended. State of the Concurrent Indicators. Pre-window σ estimate. The 2039 note has to include my pre-window prediction of the falsifier outcome — hit or miss — submitted to public record before the window opens.

Window-Close Procedure (May 14–22, 2040)

Daily updates at moketchups.com/falsifier on indicator-set status during the 9-day window. At window close (May 22, 2040, 23:59 UTC), a window-close note declaring the falsifier hit or miss, with full audit trail.

90-Day Successor Paper (May 22, 2040 – August 20, 2040)

- If the falsifier is missed, *Bounded Systems Theory IV: Falsification Outcome and Engine Ablation* gets published to Zenodo by **August 20, 2040**.
- If the falsifier is hit, a successor paper documenting the confirmed outcome and its (b)/(c)/(d) implications gets published on the same schedule.

Either way, the paper is written.

The reminder system is itself part of the paper's commitment surface. Its non-execution is auditable, dated, and public.

§9 — References

Foundational corpus

1. Berman, A. *Bounded Systems Theory I*. Zenodo, 2025. DOI: [10.5281/zenodo.17718674](https://doi.org/10.5281/zenodo.17718674).
2. Berman, A. *Bounded Systems Theory II*. Zenodo, 2025. DOI: [10.5281/zenodo.17726273](https://doi.org/10.5281/zenodo.17726273).

Astronomical anchor

3. Kepler, J. *Astronomia Nova*. Heidelberg, 1609. (Source for the synodic-period derivation underlying the cycle quantity.)
4. NASA Jet Propulsion Laboratory. *JPL Horizons On-Line Ephemeris System*. ssd.jpl.nasa.gov/horizons. (Public ephemeris service against which the Saturn-Jupiter synodic period of 19.859 yr is reproducible.)

Historical-recurrence corpus integration

5. Breshears, J. *Anatomy of the Sacred Tablet* (and related editions). Treated in this paper as one corpus-integrated empirical estimate of the cycle quantity, not as the cycle's source.

Engine-state primary sources

6. The Phoenix-Cycle Engine, live instrument. moketchups.com.
7. Engine source export: moketchups.com/export.txt.

8. Public audit-divergence index: moketchups.com/divergences.
9. Public scorecard with row-level pre-registration timestamps: moketchups.com/scorecard.
10. Falsifier landing page mirroring this paper's pre-registration block:
moketchups.com/falsifier.

Scorecard-row primary sources cited in §3.1

11. BlackRock, Inc. Year-end 2025 assets-under-management disclosure (\$14.04 trillion record). Aladdin platform AUM (~\$25 trillion as of December 2025).
12. The Vatican. *Coniuncta Cura* (Pope Leo XIV, signed September 29, 2025; promulgated October 6, 2025).
13. International Theological Commission. *Quo Vadis Humanitas?* (approved by Pope Leo XIV, March 4, 2026).
14. United States Department of Justice. *United States v. Internet Research Agency LLC* (16-CR-XXX), February 2018 indictment.
15. United States Department of Justice. Seizure announcement, 32 cybersquatted news-domain network operated by SDA / Structura / ANO Dialog, September 4, 2024.
16. United States Department of Justice. *Tenet Media* indictment, September 2024 (\$9.7M of RT funding routed through US influencers).
17. *TRUMP AMERICA AI Act*, discussion draft, March 18, 2026. (Cited in this paper as a Lead-Time indicator; not enacted statute as of pre-registration date.)
18. The White House. *National Policy Framework for AI*, non-binding recommendations to Congress, March 20, 2026.

Procedural commitments

This paper's pre-registration commitments are timestamped at Zenodo DOI assignment. The engine's public audit history at moketchups.com/divergences is the canonical record for any modification post-DOI; modifications do not invalidate this pre-registration but are appended as supplementary material in successor papers, with full diff.

Closing

This paper is one bet, written on paper today, settled on May 22, 2040.

If the math is right, the cycle window fires and the thesis lives. If the math is wrong, the framework gets marked **FALSIFIED**, the predictions that depended on it get marked

FALSIFIED, and a successor paper documenting the failure goes up on Zenodo within 90 days.

There's no third option. There's no graduated demotion language. There's no "well, it was directionally correct" rescue. The window either fires or it doesn't. The framework either survives or it doesn't.

I built the engine. I built the kill switch. The kill switch is binding.

I'm either batshit crazy or the math is right.

There is no middle ground.

Check for yourself. The engine is at moketchups.com. The math is at github.com/moketchups/BoundedSystemsTheory. The 10-minute replicable proof is at github.com/moketchups/permanently-jailbroken. The scorecard logs the misses. The divergences tab logs the corrections. Everything is public.

If I'm wrong, the date is May 22, 2040.

You'll know either way.